

DS 503: Advanced Data Analytics

Week 2: Projection Techniques

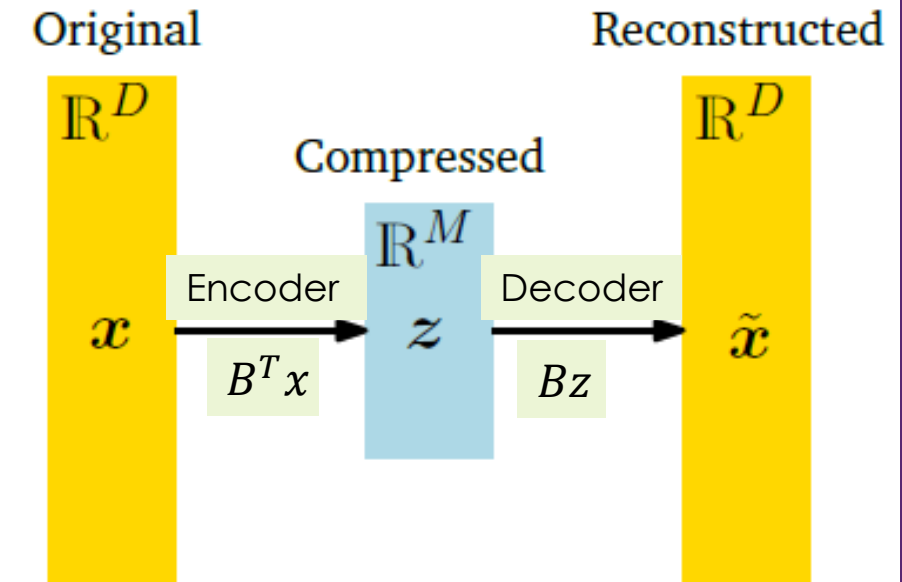
Gagan Raj Gupta

How to deal with high-dimensions?

- Project data to a low-dimensional space
- Fidelity: While doing so, don't disturb the relative distances
- Where is this used?
 - Principal component analysis
 - Nearest Neighbor Search
 - Clustering
 - Text Embeddings
 - Graph Embeddings
 - Image manifolds

Problem statement

- $D \times n$ data matrix A (D rows and n columns)
 - Each column is a D -dimensional vector
 - Columns are normalized (mean = 0)
- Assume there is a projection matrix B , such that
 - $z = B^T x \in \mathbb{R}^M$
 - $B = [b_1, b_2, \dots, b_M]$ has orthogonal columns
- B is the best-fit M -dim. subspace S_M for rows of A
 - Minimize reconstruction error
 - Minimize sum of squared distances from A_i to S_M




Minimize:

$$||x - \tilde{x}||^2$$


Best fit subspaces and Maximizing Information

- Let's begin with the following question.
 - Find a direction in which the data ($\mathbf{D} \times \mathbf{n}$ matrix) has maximum information
 - Maximize $|\mathbf{A}v_1|$, such that $|v_1| = 1$
 - This is also equivalent to Minimizing the sum of squared distances to the point nearest to the line.
- By successively solving the above problem, we can find the best-fit k -dimensional subspaces
 - Which preserve the maximum possible information (by maximizing projections)
 - Or equivalently, minimize the sum of squared distances of the vectors to the subspace
- When $k=r$, the rank of A , we get SVD



Projections, Distances and Data Variance

- Minimizing distance = maximizing projection
 - $\|x\|_2^2 = (\text{projection})^2 + (\text{distance to line})^2$
- SVD: Find best fit 1-dimensional line
 - u_1 = unit vector along the best fit line
 - x_i = i -th column of A , length of its projection: $|\langle x_i, u \rangle|$
- Sum of squared projection lengths: $\|A^T u\|_2^2$
- **First singular vector:**
$$u_1 = \arg \max_{\|u\|_2=1} \|A^T u\|_2$$
- If there are ties, break arbitrarily
 - $\sigma_1(A) = \|A^T u_1\|_2$ is the **first singular value**



Variance of z_1 of $z \in R^M$

$$V_1 := V[z_1] = \frac{1}{N} \sum_{n=1}^N z_{1n}^2$$

$$z_{1n} = b_1^T x_n$$

$$V_1 = b_1^T S b_1$$

Where S is the data-covariance matrix

$$S = \frac{1}{N} \sum_{n=1}^N x_n x_n^T.$$

Maximizing Variance using PCA

Optimization Problem

$$\max_{\mathbf{b}_1} \mathbf{b}_1^\top \mathbf{S} \mathbf{b}_1$$

$$\text{subject to } \|\mathbf{b}_1\|^2 = 1.$$

Encoder

$$\mathbf{z}_{n1} = \mathbf{b}_1^\top \mathbf{x}_n \in \mathbb{R}$$

Decoder

$$\tilde{\mathbf{x}}_n = \mathbf{b}_1 \mathbf{z}_{n1} = \mathbf{b}_1 \mathbf{b}_1^\top \mathbf{x}_n \in \mathbb{R}^D$$

Lagrangian

$$\mathcal{L}(\mathbf{b}_1, \lambda) = \mathbf{b}_1^\top \mathbf{S} \mathbf{b}_1 + \lambda_1 (1 - \mathbf{b}_1^\top \mathbf{b}_1)$$

Differentiating

$$\frac{\partial \mathcal{L}}{\partial \mathbf{b}_1} = 2\mathbf{b}_1^\top \mathbf{S} - 2\lambda_1 \mathbf{b}_1^\top, \quad \frac{\partial \mathcal{L}}{\partial \lambda_1} = 1 - \mathbf{b}_1^\top \mathbf{b}_1,$$

$$\mathbf{S} \mathbf{b}_1 = \lambda_1 \mathbf{b}_1,$$

$$\mathbf{b}_1^\top \mathbf{b}_1 = 1.$$

\mathbf{b}_1 is an eigenvector of \mathbf{S}
 λ_1 is the eigenvalue

$$V_1 = \mathbf{b}_1^\top \mathbf{S} \mathbf{b}_1 = \lambda_1 \mathbf{b}_1^\top \mathbf{b}_1 = \lambda_1,$$

- Variance of the data projected onto a one-dimensional subspace equals the **eigenvalue** that is associated with the **basis vector \mathbf{b}_1** that spans this subspace.
- To maximize the variance of the low-dimensional code, we choose the basis vector associated with the **largest eigenvalue** principal component of the **data covariance matrix**.
- This eigenvector is called the first *principal component*

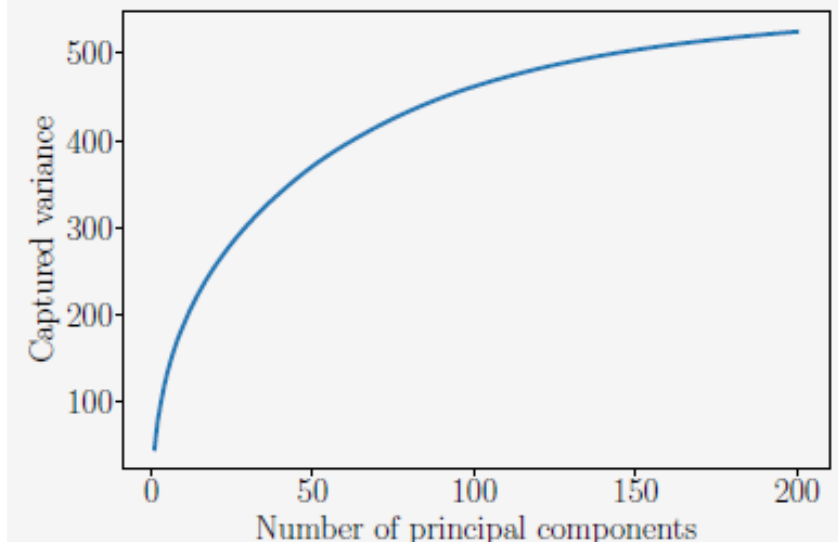

General Algorithm

- In general, the largest M (orthonormal) eigenvectors of the data-covariance matrix span the best M dimensional subspace for A
- They also capture the variance equal to the sum of largest M eigenvalues
- We can iteratively compute the largest eigenvalue/eigenvector of the covariance matrix 'S' and continue till the sum of M eigenvalues captures a large fraction of the trace of S

$$V_M = \sum_{m=1}^M \lambda_m$$

$$1 - \frac{V_M}{V_D}.$$

Stop when this
become small

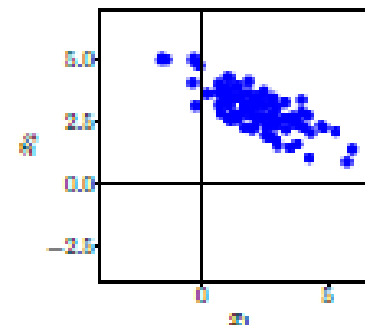


Key Steps of PCA in Practice


- Centering of the data
- Normalization
- Eigen decomposition of the Covariance Matrix
- Projection
- Mapping back to the original data space

10.6 Key Steps of PCA in Practice


337




(a) Original dataset.



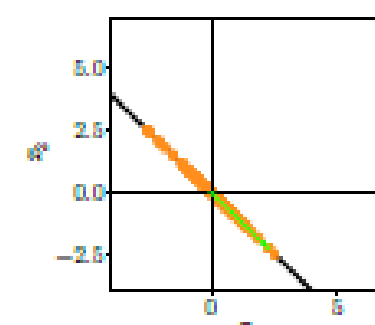
(b) Step 1: Centering by subtracting the mean from each data point.




(c) Step 2: Dividing by the standard deviation to make the data unit free. Data has variance 1 along each axis.



(d) Step 3: Compute eigenvalues and eigenvectors (arrows) of the data covariance matrix (ellipse).



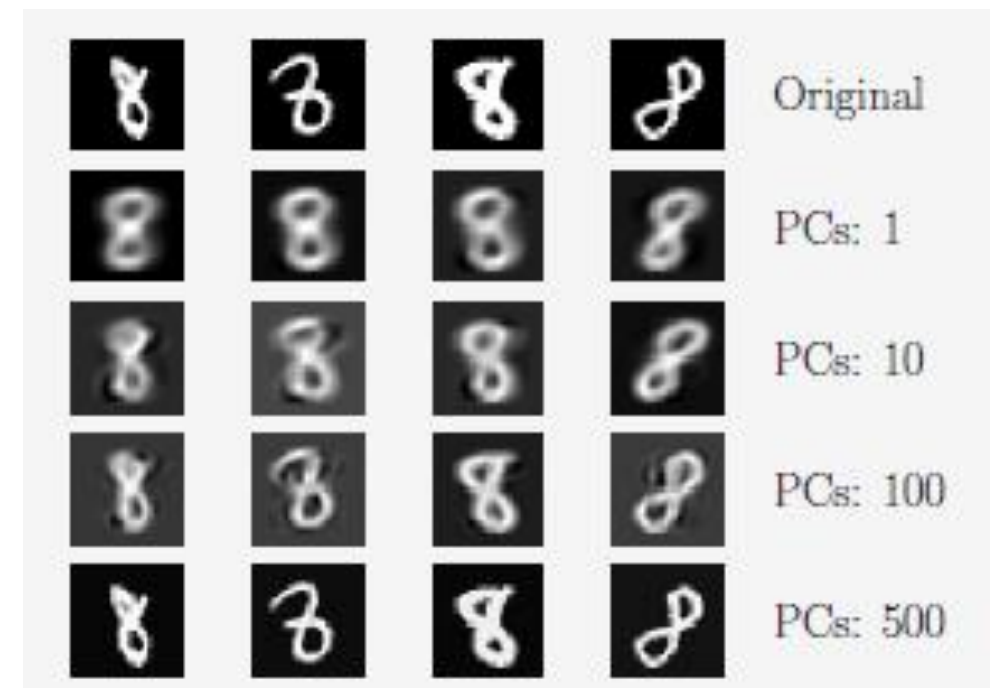
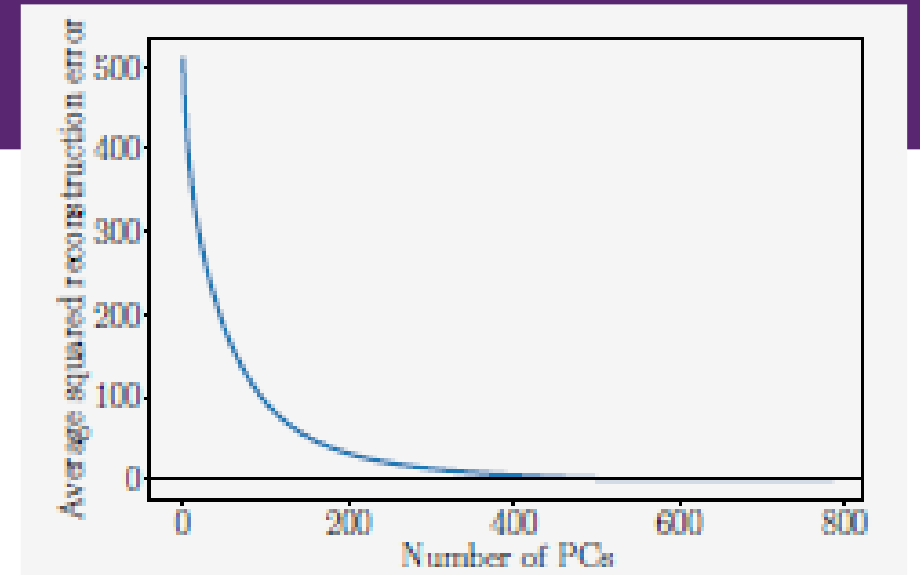
(e) Step 4: Project data onto the principal subspace.



(f) Undo the standardization and move projected data back into the original data space from (a).

MNIST Digits: reconstruction

- MNIST dataset contains of handwritten digits 0 to 9
- Each digit is an image of 28x28 pixels, that is equivalent to a vector of dimension 784
- The picture on the right is performing PCA on the 5389 samples of “8” from the dataset
- We check the reconstruction error curve to decide how many components to keep



Variants of PCA

- **Sparse PCA** : Add a L1 regularization term to remove features with small coefficients. Find sparse low rank matrices B and C so that $A \approx BC$
- **Non-negative Matrix Factorization**: Find non-negative matrices U and V so that $A \approx UV$
- **Kernel PCA** : Account for a non-linear relationship among features by using the “kernel trick”
- **Probabilistic PCA (PPCA)**: Formulates PCA as a generative process which allows us to simulate new data

Application Areas:

- **Facial Feature Extraction**

Find a few basic faces so that every face can be expressed as a combination

- **Gene Expressions**

Find out which genes are playing a role in a particular phenomenon

- **Text Mining and Document Classification**

Combine the various topics in a document

Properties of SVD

- SVD: $A = U\Sigma V^T$
- Σ = Diagonal matrix (positive real entries σ_{ii})
- U, V : orthonormal columns:
 - $\mathbf{u}_1, \dots, \mathbf{u}_r \in \mathbb{R}^D$ (directions that maximize projections of x_i)
 - $\mathbf{v}_1, \dots, \mathbf{v}_r \in \mathbb{R}^n$ (projections of x_i on u_i)
 - $\langle \mathbf{u}_i, \mathbf{u}_j \rangle = \delta_{ij}, \langle \mathbf{v}_i, \mathbf{v}_j \rangle = \delta_{ij}$
- $A = \sum_i \sigma_i \mathbf{u}_i \mathbf{v}_i^T$
- Thus, any matrix can be represented as sum of “r” rank-1 matrices

The diagram illustrates the SVD decomposition $A = U\Sigma V^T$ with dimensions for each matrix:

- A is a $(D \times n)$ matrix, represented by a large green rectangle.
- U is a $(D \times r)$ matrix, represented by a medium green rectangle.
- Σ is a $(r \times r)$ matrix, represented by a small green rectangle.
- V^T is a $(r \times n)$ matrix, represented by a small green rectangle.

The equation is shown as: $A \text{ (D x n)} = U \text{ (D x r)} \Sigma \text{ (r x r)} V^T \text{ (r x n)}$

Singular Values vs. Eigenvalues

- If A is a square matrix:

- Vector \mathbf{v} such that $A\mathbf{v} = \lambda\mathbf{v}$; is an eigenvector with eigenvalue = λ
- For symmetric real matrices, A , \mathbf{v} 's are orthonormal
 - A can be expressed as: $A = V\Sigma V^T = U\Sigma U^T$
 - V 's columns are eigenvectors of A
- Diagonal entries of Σ are eigenvalues $\lambda_1, \dots, \lambda_n$

- SVD is defined for all matrices (not just square)

- Orthogonality of singular vectors is automatic

$$A\mathbf{v}_i = \sigma_i\mathbf{u}_i \text{ and } A^T\mathbf{u}_i = \sigma_i\mathbf{v}_i \text{ (will show)}$$

$$AA^T\mathbf{u}_i = \sigma_i^2\mathbf{u}_i \Rightarrow \mathbf{u}_i\text{'s are eigenvectors of } AA^T \text{ ((N-1) times sample covariance matrix)}$$

SVD: Greedy Construction

- Find best fit 1-dimensional line, repeat r times (until projection is 0)
- **Second singular vector and value:**

$$\mathbf{u}_2 = \arg \max_{\mathbf{u} \perp \mathbf{u}_1, \|\mathbf{u}\|_2=1} \|A^T \mathbf{u}\|_2$$
$$\sigma_2(A) = \|A \mathbf{u}_2\|_2$$

- **k-th singular vector and value:**

$$\mathbf{u}_k = \arg \max_{\mathbf{u} \perp \mathbf{u}_1, \dots, \mathbf{u}_{k-1}, \|\mathbf{u}\|_2=1} \|A^T \mathbf{u}\|_2$$
$$\sigma_k(A) = \|A^T \mathbf{u}_k\|_2$$

- We can show that: $(\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_k)$ is best-fit subspace

Best low rank approximations of a Matrix

- Suppose I have to find the best “k” rank approximation A_k matrix.

$$\text{Let } A_k = \sigma_1 u_1 v_1^T + \dots + \sigma_k u_k v_k^T$$

Eckart-Young: If B has rank k then $\|A - B\| \geq \|A - A_k\|$

- This result has been proven for multiple norms
- Spectral: $\|A\|_2 = \max \frac{\|Ax\|}{\|x\|} = \sigma_1$
- Frobenius: $\|A\|_F^2 = \sigma_1^2 + \dots + \sigma_r^2 = \sum_{i,j} |a_{i,j}|^2 = \text{trace of } AA^T$
- Nuclear: $\|A_N\| = \sigma_1 + \dots + \sigma_r$ (the trace norm)

Applications of low rank approximations

- Principal component analysis (fitting a hyperplane to data)
- Model reduction in analyzing physical systems
- Fast algorithms in scientific computing
- PageRank and other spectral methods in data analysis
- Diffusion geometry and manifold learning
- Many, many more ...


Computing the SVD: Power iteration method

- Begin with a random vector x_0 that is not in the null space of $S = AA^T$
- Follow the iteration $x_{k+1} = \frac{S x_k}{||S x_k||}$
- This converges to the eigenvector associated with the largest eigenvalue of S
- Also used in the page-rank algorithm for ranking web-pages based on their hyperlinks
- Issues: S can be very large
- Alternative 1: If $N \ll D$, then we can instead work with the matrix $A^T A$ which also has the same eigenvalues (square of singular values).
- This can happen, for example, when we are dealing with large size images with millions of pixels.

SVD works well for small matrices ($m, n < 5000$) or sparse matrices. The function for computing SVD is not so easy to write.

JL Lemma and Random Projections

- Suppose x_1, x_2, \dots, x_n are any n points in R^D and $k \geq \frac{8 \ln n}{\epsilon^2}$.
- \exists a distance preserving projection J from R^D to R^k which works \forall pairwise distances:
 - $$(1 - \epsilon) \|x_i - x_j\|^2 \leq \|Jx_i - Jx_j\|^2 \leq (1 + \epsilon) \|x_i - x_j\|^2$$
- There are multiple proofs and one of the proof shows that w.h.p. a random projection ($k \times D$) is very likely to keep the n points apart.
- Project each $x \in A$ onto $f(x)$, where $f: \mathbb{R}^D \rightarrow \mathbb{R}^k$
- Pick k vectors u_1, \dots, u_k i.i.d: $u_i \sim N_D(0^D, 1)$
$$f(v) = (\langle u_1, v \rangle, \dots, \langle u_k, v \rangle)$$
 - Since the k vectors were Gaussian, the projections are also Gaussian.
 - Application of Gaussian Annulus theorem allows to bound the deviation in the distances



Applications

- Suppose x_1, x_2, \dots, x_n are any n points in R^D , where D is very large. Tasks:
- Suppose the points almost live on a **linear subspace** of (small) dimension k .
 - Find a basis for the “best” subspace. (Principal component analysis.)
- Given k , find the subset of k vectors with **maximal spanning volume**.
- Suppose the points almost live on a low-dimensional **nonlinear manifold**.
Find a parameterization of the manifold.
- Given k , find for each vector x_i , its **k closest neighbors**.
- Partition the points into **clusters**.

Note: Some of these don't have well-defined solutions and some are combinatorially hard. If we can embed the points in a low-dimensional subspace, while preserving distance approximately, then a variety of algorithms for solving these problems become possible.

Application: Nearest Neighbors Search

Goal: Given a database of n points in R^D where n and D are usually large. Query points in R^D . Queries should be fast.

- If the database has n_1 points and n_2 queries are expected during the lifetime of the algorithm, take $n = n_1 + n_2$
- Project database (using random vectors) to a k -dimensional space
- Store projections in an efficient data structure (more in next lecture)
- On receiving a query, project the query to the same subspace and compute nearby database points.
- The JL Lemma says that with high probability this will yield the right answer whatever the query

Linear Regression (Problem setup)

- We are given observations (\mathbf{x}, y)
 - Let us build a linear model to predict y from the observations \mathbf{x} s. t. $w_0 + \sum_{i=1}^d w_i x_i = y$
- Compute the “best” solution to the model $X\mathbf{w} = y$ that minimizes SSE (sum of squared distances)
- This problem is very similar to solving $A\mathbf{x} = \mathbf{b}$, a system of linear equations
- In most cases, there may be no exact solution to this problem
- But, there are many approaches we can take to get the “least squares” solution
 - This minimizes the sum of square errors $\|\mathbf{b} - A\hat{\mathbf{x}}\|^2$

Normal Equations


Since $(\mathbf{b} - A\hat{\mathbf{x}})$ is perpendicular to all vectors $A\mathbf{x}$ in the column space, $(A\mathbf{x})^T(\mathbf{b} - A\hat{\mathbf{x}}) = \mathbf{x}^T A^T(\mathbf{b} - A\hat{\mathbf{x}}) = 0$

Normal Equation for solving $\hat{\mathbf{x}}$: $A^T A\hat{\mathbf{x}} = A^T \mathbf{b}$

Least squares sol to $A\mathbf{x}=\mathbf{b}$: $\hat{\mathbf{x}} = (A^T A)^{-1} A^T \mathbf{b}$

Projection of \mathbf{b} onto $\text{Col}(\mathbf{A})$: $\mathbf{p} = A\hat{\mathbf{x}} = A(A^T A)^{-1} A^T \mathbf{b}$

Projection matrix that multiplies \mathbf{b} to give \mathbf{p} : $\mathbf{P} = A(A^T A)^{-1} A^T$



Pseudo Inverse Method

○ Pseudo Inverse Method: $\hat{x} = A^+ b$

○ If A has independent columns, $A^+ = (A^T A)^{-1} A^T$


○ If A has independent rows, $A^+ = A^T (A A^T)^{-1}$

○ Pseudo inverse can be computed using SVD: $A^+ = V \Sigma^+ U^T$

○ Σ^+ contains the inverse of all non-zero diagonal elements in Σ

$$\Sigma = \begin{bmatrix} \sigma_1 & 0 & 0 & 0 \\ 0 & \sigma_2 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

$$\Sigma^+ = \begin{bmatrix} \frac{1}{\sigma_1} & 0 & 0 & 0 \\ 0 & \frac{1}{\sigma_2} & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$



QR (Gram-Schmidt) Method

- Decompose $A = QR$
 - Where Q is an orthogonal matrix, and R is a triangular matrix
- Then, $A^T A = R^T Q^T QR = R^T R$ and the normal equation $A^T A \hat{x} = A^T b$, can be solved as
 - $\hat{x} = (R^T R)^{-1} R^T Q^T b = R^{-1} Q^T b$
- This is computationally efficient to solve

Alternatives to compute Low Rank Approximations

- Let A be an $m \times n$ matrix of low numerical rank
- Suppose that you can't afford to compute the full SVD or you don't have a good implementation
- How can you compute a low-rank approximation to A ?
 - Gram-Schmidt: Keep reducing a rank-1 component from A
 - Complexity: $O(mnk)$
 - Krylov Methods: Restrict the matrix A to the k -dimensional "Krylov subspace"
 - Span $(r, Ar, A^2r, \dots, A^{k-1}r)$
 - Compute eigen-decomposition of resulting matrix

```
(1) for  $k = 1, 2, 3, \dots$ 
(2)   Let  $i$  denote the index of the largest column of  $A$ .
(3)   Set  $q_k = \frac{A(:,i)}{\|A(:,i)\|}$ .
(4)    $A = A - q_k(q_k^* A)$ 
(5)   if  $\|A\| \leq \varepsilon$  then break
(6) end while
(7)  $Q = [q_1 \ q_2 \ \dots \ q_k]$ .
```

Each of these approximations result in a factorization of the form

$$A \approx Q Q^T A$$

Where Q is an approximate orthonormal basis for the column space of A

Randomized Low Rank Approximations

Range Finding (Basis) Problem: Given an $m \times n$ matrix A and an integer $k < \min(m,n)$. Find an orthonormal $m \times k$ matrix Q such that $A \approx Q Q^T A$

Solving the primitive problem via randomized sampling — intuition:

1. Draw **Gaussian random vectors** $g_1, g_2, \dots, g_k \in R^n$
2. Form “**sample**” vectors $y_1 = Ag_1, y_2 = Ag_2, \dots, y_k = Ag_k \in R^m$
3. Form **orthonormal vectors** $q_1, q_2, \dots, q_k \in R^m$ such that

$$\text{Span}(q_1, q_2, \dots, q_k) = \text{Span}(y_1, y_2, \dots, y_k)$$

For instance, Gram-Schmidt can be used — pivoting is rarely required.

If A has exact rank k , then $\text{Span}\{q_j\}_{j=1}^k = \text{Range}(A)$ with probability 1.

Randomized SVD

- **Goal:** Given an $m \times n$ matrix \mathbf{A} , compute an approximate rank- k SVD $\boxed{A \approx U\Sigma V^T}$
- **Algorithm:**
- 1. Draw an $n \times k$ Gaussian random matrix \mathbf{G} . $G = \text{randn}(n,k)$
- 2. Form the $m \times k$ sample matrix $\mathbf{Y} = \mathbf{A}\mathbf{G}$ $Y = A * G$
- 3. Form an $m \times k$ orthonormal matrix \mathbf{Q} such that $\mathbf{Y} = \mathbf{Q}\mathbf{R}$ $[Q, \sim] = \text{qr}(Y)$
- 4. Form the $k \times n$ matrix $\mathbf{B} = \mathbf{Q}^T \mathbf{A}$ $B = Q' * A$
- 5. Compute the SVD of the small matrix \mathbf{B} : $\mathbf{B} = \hat{\mathbf{U}}\Sigma\mathbf{V}^T$ $[Uhat, Sigma, V] = \text{svd}(B,0)$
- 6. Form the matrix $\mathbf{U} = \mathbf{Q}\hat{\mathbf{U}}$ $U = Q * Uhat$
- **Power iteration to improve the accuracy:** The computed factorization is close to optimally accurate when the singular values of \mathbf{A} decay rapidly. When they do not, a small amount of power iteration should be incorporated, i.e. replace Step 2 by $\mathbf{Y} = (\mathbf{A}\mathbf{A}^T)^t \mathbf{A}\mathbf{G}$ for a small integer t , say $t = 1$ or $t = 2$.

Randomized Embeddings: “Fast” JL Transforms

- So far, the only randomized embedding we have described takes the form

$$\boxed{f: R^D \rightarrow R^k: x \rightarrow \frac{1}{\sqrt{k}} Gx}$$

where \mathbf{G} is a matrix drawn from a Gaussian distribution. Evaluating $f(\mathbf{x})$ costs $O(nk)$.

- The cost can be reduced to $O(n \log k)$ or even less, by using *structured* random maps.
 - Subsampled Fourier transforms. (Or Hadamard transform / cosine transform / . . .)
 - Sparse random embeddings — pick matrix that consists mostly of zeros.
 - Chains of random Givens' rotations.